

Infinite Software

Machine labor makes software unlimited. The rails it runs on are not.

SKYCATCHER · JULY 16, 2026 · FULL TEXT

PART ONE

I

Infinite Software

the era — machine labor makes software unlimited

“The Internet is a tidal wave. It changes the rules.”

BILL GATES · MAY 1995

INFINITE SOFTWARE · PART I

In May 1995 Bill Gates told Microsoft that the internet was a tidal wave that changes the rules, and the next thirty years of value creation proved him right. The memo is worth studying less for its accuracy than for its method.

Gates could not name Google, the smartphone, or software as a service. He did not attempt to. What he got right was magnitude, and magnitude was sufficient — Microsoft pivoted the entire company within months. Nearly every trillion-dollar business created since traces back to that wave.


We think we are standing at the same kind of moment, and that this time the wave is denominated in tokens. The rest of this piece is the evidence.

**“The Internet is a tidal wave.
It changes the rules.”**

— BILL GATES, INTERNAL MICROSOFT MEMO · MAY 26, 1995

Thirty-one years later, the wave is running again — denominated in tokens.

It marked the moment an incumbent recognized a paradigm shift. The framing was right about magnitude, even when details were fuzzy. It timestamps the beginning of a ~30-year value creation wave.



By the internet clock, it is 1996

Line the two waves up on the same clock and the analogy becomes specific rather than rhetorical.

The 1995 memo corresponds to November 2022, when ChatGPT arrived from approximately zero — the Netscape moment. From 1996 to 1998 the rails were built and paid for while the street debated whether the whole thing was a fad. That is exactly where we are: 2026, the reprice beginning, multiples still at fade levels.

What follows in the internet analogue is 1999 and 2000 — the mania, 65 to 100x, and then a compression of 60 to 85%. We are not underwriting the mania. We are underwriting 2027 and 2028, when agent disclosures arrive in earnings and the multiple re-rates, and 2029, when the meters bill at all-time highs.

By the internet clock it is 1996. After the memo, before the mania. Early enough to matter, late enough to have receipts.

Figure 1. The internet wave vs. the token wave — same clock, new denomination

THE INTERNET WAVE

1995

The “tidal wave” memo names the discontinuity

1996–98

The rails get built and paid — the street debates “fad”

1999–2000

The mania — 65–100x, then 60–85% compression

2001+

The rails' owners collect the era's fortunes

THE TOKEN WAVE

NOV 2022

ChatGPT — the Netscape moment, from ~zero

2026

The reprice begins — multiples still at fade levels

WE ARE HERE

2027–2028

Agent GAAP disclosures, earnings accelerate and multiples re-rate

2029+

Meters hit all-time highs, software infra at premium multiples

Historical analogy illustrative; internet-era multiples per the acceleration-episodes analysis later in this deck.

Executive Summary

The whole argument, in three phrases. The era is Infinite Software: machine labor makes producing software effectively free, so the population of software explodes. The phenomenon is the Token Tidal Wave: every price drop buys more demand than it gives up, roughly 27x tokens by 2029. The bet is a Royalty on Machine Labor: infinite software runs on finite, metered rails, so own the meters and collect the royalty.

Everything after this page is evidence for those three lines.

Figure 2. The argument of this presentation

THE ERA	THE PHENOMENON	THE BET
Infinite Software Machine labor makes software free — the software population explodes.	The Token Tidal Wave Every price drop buys more demand than it gives up — ~27x tokens by 2029.	A Royalty on Machine Labor Infinite software runs on metered rails. Own the meters; collect the royalty.

AI is the biggest tech wave yet — and this one produces labor

The institutional investor sizes artificial intelligence against the software budget. It is the natural comparison, because every technology wave of the last sixty years sold software or the hardware to run it, and each one was underwritten that way. We think the comparison is the single largest analytical error in the market today.

Mainframes, the personal computer, networking, the web, mobile, cloud — every one of them sold a tool that made a person more productive, and every one of them minted the infrastructure names of its era. IBM and DEC. Microsoft and Intel. Cisco. Google and Amazon. Apple. AWS.

This wave does not sell a tool. It sells the labor itself. That is a categorical difference, not a matter of degree, and it changes the denominator: the addressable market is not what companies spend on software, it is what they spend on work. The infrastructure names of this era are being minted right now, and they are being priced as though the old denominator still applies.

Figure 3. Skycatcher illustrative view of tech wave scale

1960–80	1980S	1990S	2000S	2010S	2015–20	2022–2029+
Mainframe	PC	Networking	Web 1.0	Mobile	Cloud / SaaS	Artificial Intelligence
IBM DEC	Microsoft Intel	Cisco Nokia	Google Amazon	Apple Tencent	AWS Salesforce	NVIDIA OpenAI Anthropic

Skycatcher illustrative framework. Company references informational only.


Software production just broke free of human hands

The cost of producing software has fallen for thirty years, and until recently it fell politely — one step roughly every decade, each step large enough to reorganise the industry and small enough to absorb. Artisan engineering ran near \$200 a unit of work. Offshoring took it to \$60. Open source and reuse took it to \$20. Copilots took it to \$5.

Agent production took it to five cents, and it did so in a single step. That is roughly 4,000x, arriving inside one product cycle rather than one decade.

We would ask you to hold on that number rather than move past it, because the magnitude is the whole argument. When any input becomes four thousand times cheaper, no one buys the same quantity of it more cheaply. They buy vastly more of it, and they buy it for uses that were previously unthinkable. Every forecast that treats this as a margin story has already missed the point.

Figure 4. Cost per unit of software work, by production regime (\$, log scale, illustrative)



Skycatcher illustration; inference economics per Epoch AI constant-performance price data. For illustrative purposes only.

Labor is no longer human-limited

The price collapse is the symptom. The regime change underneath it is the cause, and it is simpler than the discourse suggests. A copilot keeps a person inside every loop: you prompt, it answers, you check, you prompt again. An agent takes an objective and runs — planning, acting, verifying, looping until the work is finished. The human stops being the operator and becomes the manager.

Three consequences follow, and they multiply rather than add. The human bottleneck disappears, so attention-hours no longer cap output. Work parallelises, because agents launch agents — ten to a hundred simultaneous workstreams from one instruction. And the duty cycle changes: forty attention-hours a week becomes one hundred and sixty-eight machine-hours.

One analyst helping you becomes a thousand analysts working around the clock. We do not claim this substitution is clean today; reliability and supervision remain real constraints, and anyone who has run agents at scale knows it. We claim the direction is not in question.

Figure 5. From copilots that help you, to agents that work for you

Copilots / Chat

1 analyst helping you

- You prompt
- It answers
- You check
- You prompt again...

Human required in every loop

Agents

1,000 analysts working 24/7

- You define the objective
- Agents plan, act, verify
- Tool calls · writes · tests
- Loop until completion

Human becomes the manager

Key insights

- Removes the human bottleneck — attention-hours no longer cap output
- Parallelization — agents launch agents, 10–100x workstreams
- Always-on duty cycle — 40 attention-hours → 168 machine-hours a week

Skycatcher framework. Agent capabilities illustrative; reliability and supervision requirements remain constraints.

Making software is becoming free. Running it never is.

Now follow the money through the collapse, because this is where we part company with consensus most sharply.

Everything on the left-hand side of the ledger deflates toward zero: code, tests, documentation, integration glue, one-off tools, and eventually the applications themselves. That side becomes an infinite knife-fight of entrants, each of them able to produce what used to take a funded team, none of them able to hold a price. It is the worst place in the value chain to own an asset.

Everything on the right-hand side is metered forever. Compute cycles and runtime. State, memory, and context. Delivery and network. Identity, audit, and the billing rails themselves. None of it is free, none of it deflates to nothing, and all of it scales with the population of software rather than the price of writing it.

The economic framing here is not new. Jevons described it in 1865: make an input more efficient and its consumption expands rather than shrinks. Nadella invoked exactly this in January 2025. An economist will fairly object that the modern literature calls this the rebound effect, and that full backfire — consumption growing faster than efficiency improves — is historically rare, because it requires demand elasticity above one. That is precisely the point: we are not assuming the rare case, we have measured it, four years running, and Part II shows the slope. Where elasticity sits is where this thesis lives or dies, which is why it leads our kill conditions.

One more objection belongs here, because every IT professional raises it within a minute: if agents make building software free, don't they also make running your own rails free — why pay a meter when an agent can babysit your own Postgres? Partly, yes, and we hold the layers most exposed to that lightly. But identity, audit, and compliance cannot be self-hosted away, because their product is not uptime — it is a third party to blame, and no enterprise lets an agent be the counterparty of record. We return to this at full length in the conclusion.

Figure 6. The scarcity inversion — where the money goes when production costs collapse

Deflates to zero

- Code, tests, documentation
- Integration glue & boilerplate
- One-off tools & scripts
- Custom applications themselves

an infinite knife-fight of entrants

Metered forever

- Compute cycles & runtime
- State, memory & context
- Delivery & network
- Identity, audit & billing rails

billed per unit of use — scales with the software population


Last time production costs collapsed, the rails got paid first

We have a precedent for this, and it is close enough to be uncomfortable. The last time the cost of production collapsed, the rails were paid first and they were paid enormously — before the mania, while the argument was still being had in public.

In the four years following the tidal wave memo, from May 1995 to May 1999, the NASDAQ roughly tripled. Intel returned about 6x. Microsoft about 7x. Cisco — which sold nothing to consumers and everything to the people building the network — returned about 13x.

Every one of those returns was earned before 1999 began. The crowd spent those years debating whether the internet was a fad. The rails simply compounded. Past performance is not a forecast, and we are not offering these numbers as one. We are pointing out that the shape of the opportunity in front of us has been observed before, and that the observation was available to anyone willing to underwrite the infrastructure while the narrative was still contested.

Figure 7. Total return, 4 years following the internet tidal wave memo (May 1995 – May 1999)



Approximate total-return multiples from public market data, May 1995–May 1999. Past performance is not indicative of future results.

PART TWO

II

The Token Tidal Wave

the phenomenon — our ~27x call, landing on the rails

*“Efficiency in the use of an input expands, rather than reduces,
its consumption.”*

W. S. JEVONS · 1865

INFINITE SOFTWARE · PART II


Tokens grew ~2,000x in four years

Before any forecast, the record. Global token consumption went from 0.05 quadrillion a year in 2022 to roughly 100 quadrillion in 2026. That is about 2,000x in four years, and it arrived in an orderly sequence: 10x, then 10x, then 5x, then 4x.

Note what that line ran through. Two market corrections and one full-scale AI panic, complete with declarations that the buildout was a bubble and the models had hit a wall. Demand never blinked. Not once did the curve bend to accommodate the narrative.

These are aggregated platform disclosures, not a survey and not our imagination. The honest caveats: tokenizers differ across labs, so cross-platform token counts are not perfectly commensurable, and the disclosure anchors are sparse — a “100 trillion a quarter” here, a “quadrillion a month” there — so the series carries error bars of tens of percent. None of that matters to the argument. The uncertainty is in the second digit; the phenomenon is in the exponent.

Figure 1. Global token consumption, quadrillion tokens / year (log scale, Skycatcher aggregation of platform disclosures)



Platform disclosures (e.g., 100T tokens/quarter, 1Q tokens/month anchors), Skycatcher aggregation; error bars apply. Illustrative.

Every 10x price drop has bought 15–20x more demand

The institutional model of this market assumes that falling prices compress revenue. It is the reflex of anyone who has watched a hardware cycle. Applied here it produces the fade case that dominates every sell-side model, and it has been wrong for four consecutive years.


What actually happened: frontier price per million tokens fell about 50x, from roughly \$20 to \$0.40. Volume rose about 2,000x. Every 10x reduction in price has purchased 15 to 20x more demand — a measured elasticity of about 1.6, sustained across the entire period.

We checked the obvious objection, which is that buyers were purchasing capability rather than responding to price — the confound would be that each year's tokens simply do more. So we ran the same measurement against constant-capability tiers: models held at a fixed benchmark level whose prices fell as newer tiers arrived above them. The slope holds. Cheaper tokens do not shrink this market. They enlarge it, and they enlarge it faster than they cut it.

A note on terms, because we use two elasticities in this document and they are not the same number. The 1.6 here is demand elasticity — tokens bought per unit of price cut. The 0.5 to 0.7 in Parts IV and V is monetization elasticity — revenue captured per unit of usage growth, after discounts and commitments. The first says the wave builds; the second, deliberately conservative, says how much of it the meters keep.

Figure 2. What happened, 2022–26 — price down ÷50x, volume up ×2,000x, on the measured slope

LOG SCALE · PRICE FALLING AGAINST VOLUME RISING



The heaviest users barely exist yet

The reason we think this continues is that the heaviest users of tokens barely exist yet.

Of roughly a billion weekly AI users, about 2.5% run agents today. That alone is 40x of headroom in penetration. But penetration is the smaller half of the arithmetic. A casual chat user consumes on the order of 10,000 tokens a day. An always-on agent consumes 5 to 25 million — a difference of 500 to 2,500x for the same human being.

And the clock is different. A person gives you about forty attention-hours a week. An agent runs one hundred and sixty-eight. The same customer becomes worth orders of magnitude more the day their software stops waiting to be prompted.

Figure 3. Where the 27x comes from — penetration × intensity × price

2.5%

of ~1B weekly AI users run
agents today

→ 40x headroom in users alone

500–2,500x

tokens per user-day

casual chat ~10K vs. always-on agent
5–25M

168 hrs

agent duty cycle per week

vs. ~40 human attention-hours

Skycatcher estimates from platform usage disclosures; intensity band observed in the wild, illustrative.

Tokens are shifting from “asking” to “doing”


This is the assumption that matters most in the entire piece, and we would rather name it plainly than bury it. Tokens are shifting from asking to doing.

Today roughly 85% of tokens are spent on chat and 15% on delegated work. Our case requires that ratio to invert to something like 25/75 by 2029. If the delegation ladder stalls, our numbers do not happen.


The evidence that it is not stalling is already public. About 77% of one major lab’s API traffic already classifies as automation rather than conversation. Coding’s share of routed model traffic went from 11% to over 50% in a single year. Enterprise API tokens at the leading lab rose 150% in five months. Our case needs a doing-share of 60% or better; anything above that beats consensus outright.

Figure 4. Share of tokens by mode — the delegation ladder

TODAY



2029 — OUR CASE



Lab and marketplace traffic disclosures, 2025–26. Our case needs doing-share ≥60% by 2029; anything above that already beats consensus.

Our call: ~2,700Q tokens by 2029 — a 27x from here


So here is our call, stated as a number that can be graded against us: roughly 2,700 quadrillion tokens a year by 2029. From today's 100 quadrillion, that is 27x.

It sounds aggressive until you look at what it assumes about growth. Our case runs at about 3x a year. The trailing four years ran at about 7x a year. We are underwriting a deceleration of more than half, and we still arrive at a number two and a half times the street's floor of roughly 1,100 quadrillion.

The floor is not a bear case invented to look balanced. It is what you get by applying unit elasticity to the announced price path — the street's own arithmetic, run honestly. The distance between the two lines is the entire opportunity.

Figure 5. Quadrillion tokens / year (log scale) — tidal-wave case vs. the street's floor

QUADRILLION TOKENS / YEAR · LOG SCALE



Skycatcher scenarios. Floor applies unit elasticity to the announced price path; call applies the measured slope at roughly half its trailing pace.

2,700Q, decomposed: three assumptions, multiplied

The 27x is not a curve we drew and then justified. It is three assumptions multiplied, each of which can be argued with separately.

Weekly AI users roughly double, from about 1 billion to about 2 billion. The share running always-on agents goes from 2.5% to about 15%, a factor of six. And each agent-day consumes somewhere in the observed band of 10 to 25 million tokens. Two billion users, times fifteen percent, times twenty-five million tokens a day, times three hundred and sixty-five days, is approximately 2,700 quadrillion.

Notice that the floor and the wave share the same users and the same penetration. At 10 million tokens a day you get the street’s 1,100 quadrillion. At 25 million you get our 2,700. The entire debate between us and consensus is a debate about intensity — how hard a delegated agent works — and that is a question the filings will settle within four quarters.

Figure 6. The 2029 end state, bottom-up — each factor against today

~2B		~15%		10–25M		~2,700Q
weekly AI users	×	run always-on agents	×	tokens per agent-day	=	tokens per year
×2 vs today		×6 vs today		the observed band		×27 vs today's 100Q
~1B today → ~2B 2029e		2.5% today → 15% 2029e		10M floor → 25M wave		

Skycatcher estimates, illustrative: 2B × 15% × 25M tokens/day × 365 ≈ 2,700Q; at 10M/day ≈ 1,100Q (the floor). Equivalent to ×3/yr for three years vs. the trailing ~×7/yr.

Tokens are the fuel. Infrastructure events are the bill.

Here is where our thesis departs from the one you have already heard. Tokens are the fuel. They are not the bill.

Take a single delegated task — “fix the failing build.” The agent thinks, consuming perhaps 200,000 tokens of planning, reading, and reasoning. Then the agent works, and working means API calls, database reads and writes, test runs, builds, traces, and logs. One multi-step coding task can generate more than 100,000 metered events.

Two numbers here need to be squared, because a careful reader will divide one by the other. The 100,000-plus figure counts raw metered events — every log line, trace span, and cache read the task touches. Most of those are bundled or priced near zero. When we weight events by what actually bills, the same task nets out to 20 to 40 billable-weighted events per thousand agentic tokens, and that is the conservative number every downstream forecast in this document uses. Machine labor does not arrive as a token invoice. It arrives as an invoice on somebody else’s rails.

Figure 7. What one delegated task does to the meters

One delegated task

“fix the failing build”

100,000+

metered events from a single multi-step coding task

Agent thinks

~200K tokens of planning, reading, reasoning

Agent works

API calls · DB reads & writes · test runs · builds · traces · logs

20–40

billable events dragged behind every 1,000 agentic tokens

Task-level traces, Skycatcher agent stack and public platform documentation. Illustrative.

The wave doesn't stop at tokens — it lands on the meters

This is the second-order effect that consensus has not modelled at all, because consensus is still reading this as a story about the labs.

Token demand grows about 27x, past two quintillion tokens a year by 2029. That is the labs' story and it is the one being told. But every action an agent takes lands on rented rails — something that runs it, something that stores its state and memory, something that secures its identity and audits its supply chain, something that observes its logs and traces.

Run the attach rate through those four layers and billable infrastructure events grow about 130x, from roughly 0.3 quadrillion to about 39 quadrillion. The wave does not dissipate as it lands. It amplifies.

Figure 8. First-order explosion in tokens; second-order explosion in metered events

Token demand	Every action lands on rented rails	Billable infra events
~27x past two quintillion tokens a year by 2029 the labs' story	- RUNS — runtime, compute & edge - STORES — data, state & memory - SECURES — identity, audit & supply chain - OBSERVES — logs, traces & telemetry	~130x 0.3Q → ~39Q by 2029 — the wave amplifies as it lands the meters' story

Attach of 20–40 events per 1K agentic tokens per the conversion model; ~130x = 0.3Q → ~39Q billable-weighted events by 2029, Skycatcher case. Illustrative.

The wave amplifies as it lands: events grow ~130x


The amplification is not a one-off step; it compounds annually, and it compounds faster than the thing causing it.

Indexed to 2026, infrastructure events run 6x by 2027, 25x by 2028, and 130x by 2029 — between four and six times a year. Tokens over the same period run 3x, 9x, and 27x. Events outgrow tokens in every single year of the forecast, because two things deepen simultaneously: the agentic share of tokens, and the number of billable events each agentic token drags behind it.

For the labs this is a token story. For the meters it is an event story. We would rather own the event.

Figure 9. Growth index, 2026 = 1 (log scale) — billable-weighted infrastructure events vs. tokens

GROWTH INDEX, 2026 = 1 · LOG SCALE · DASHED = MODELLED



Skycatcher conversion model. For the labs this is a token story; for the meters it is an event story — events compound faster.

The tidal wave, in three numbers

Three numbers carry this part, and each is graded quarterly against public filings through Sky1.

Tokens 27x by 2029 — 3x a year against a 7x history, which makes our call a slowdown rather than a stretch. About 75% of tokens doing rather than asking, where anything above 60% already beats consensus. And roughly 130x billable infrastructure events, because the wave amplifies as it lands on the meters.

The middle number is the one to argue with. It is the most important assumption in the machine labor thesis, and everything downstream of it — revenue, margin, multiple — is a consequence rather than a separate bet. What Part III does is show that the consequences have already started appearing in the filings.

Figure 10. Part 02 recap — and what part 03 must prove

~27x

tokens by 2029

3x a year vs. a 7x/yr history — the call is a slowdown

~75%

of tokens doing, not asking

the delegation ladder — above 60% already beats consensus

~130x

billable infrastructure events

the wave amplifies as it lands on the meters

Every figure on this page is graded quarterly against public filings via Sky1.

PART THREE

III

The Receipts

The evidence in the reports and filings

“Nobody prices a thing that isn’t selling.”

SKYCATCHER

INFINITE SOFTWARE · PART III

Machine billable units have arrived — 7 major companies so far

Machine labor now has an invoice line, and that is a more important fact than any forecast in this document. Between 2023 and 2026, seven vendors shipped a machine billable unit — a price on work an agent performs rather than on a human seat. All of them dated. All of them public.

Intercom priced the outcome first, at 99 cents per resolution in 2023. Cognition priced the agent’s time, roughly \$2 per agent compute unit — fifteen minutes of work. Zendesk took \$1.50 per automated resolution in August 2024. Salesforce went to a dime an action in October 2024, down from \$2 a conversation, and has since disclosed 3.8 billion agentic work units. Cloudflare priced machine reading itself with Pay Per Crawl in July 2025. Microsoft metered Copilot Credits at a cent apiece through Azure in September 2025. ServiceNow shipped assists per agentic action in April 2026.

Four flavours emerge — outcomes, actions, sessions, metered compute — and the variety matters less than the fact of it. When an industry invents a unit of account, the market has arrived. Nobody prices a thing that isn’t selling.

Figure 1. The invoice line for machine labor — seven vendors shipped one, 2023 → 2026 · all dated, all public

2023	Intercom Fin · \$0.99 per resolution — the first outcome price
2024	Cognition Devin · ~\$2 per ACU — 15 min of agent work
AUG 2024	Zendesk · ~\$1.50 per automated resolution
OCT 2024	Salesforce · \$0.10 per action (from \$2/conv) — 3.8B AWUs
JUL 2025	Cloudflare · \$ / crawl Pay Per Crawl — machine reading, priced
SEP 2025	Microsoft · \$0.01

APR 2026 **ServiceNow • 25–150**
assists per task / agentic action

Company pricing pages and announcements, 2023 – Apr 2026. Salesforce AWU is the disclosed volume metric; billing runs on conversations and per-action credits.


The tape has turned: retention is compounding again

Retention is where a usage business tells the truth about itself, and the tape has turned.

Aggregate net dollar retention across our five public meters has risen for six consecutive quarters: 113% in the first quarter of 2025, then 113, 114, 115, 117, and 120% by the second quarter of 2026. Six quarters up, without interruption.

The reason is structural rather than commercial. Seat-based software retains at roughly 101% — you grow by selling more chairs. Consumption platforms retain at 120 to 125%. AI agent products are retaining at roughly 132%. Every task an agent runs meters more revenue from a customer who has already signed, with no salesperson involved in the increment. This is what a royalty looks like before anyone calls it one.

Figure 2. Aggregate net dollar retention, five public meters (% , Skycatcher aggregation)



Company filings, Skycatcher aggregation (approximate). Every task an agent runs meters more revenue from the same customer — no salesperson involved.

Contracted backlog is compounding at +50% against +23% revenue


Retention tells you what happened. Contracted backlog tells you what is about to.

Aggregate remaining performance obligations across the five meters went from \$10.1 billion in the first quarter of 2025 to \$19 billion in the second quarter of 2026 — roughly +50% year on year against revenue growth of about +23%. Read the sequential prints rather than the annual: +13%, +11%, +13%, +15%, +16%. The growth rate is itself accelerating.

Under committed-usage contracts, backlog leads revenue by one to three quarters. Customers are signing for consumption they have not yet drawn down. The demand is already contracted; it simply has not been recognised.

Figure 3. Aggregate remaining performance obligations, five public meters (\$B) — with q/q growth

RPO \$B



Company filings, Skycatcher aggregation (approximate). RPO leads revenue by 1–3 quarters under committed-usage contracts.

Revenue overage is accelerating


If you want the purest read on agent usage available from outside a company, it is overage — usage billed above committed floors, at list price.

We estimate it went from \$0.55 billion a quarter in the first quarter of 2025 to \$1.17 billion in the second quarter of 2026, with sequential growth of +13%, +15%, +17%, +18%, +19%. As a share of aggregate revenue it moved from roughly 14% to roughly 24% in six quarters, and it gets faster every quarter.

Overage is spending nobody budgeted. It is what happens when a workload outruns its own contract, and where limits are enforced, more than 75% of customers who hit them keep paying anyway. This is our own estimate built from on-demand revenue, consumption-versus-capacity timing, and credit disclosures — the softest number in this section, and we flag it as such.

Figure 4. Usage billed above committed floors, at list — estimated, five public meters (\$B / quarter)

OVERAGE \$B / QTR



Skycatcher estimate from on-demand revenue, consumption-vs-capacity timing, and credit disclosures. Where limits are enforced, >75% who hit them keep paying (Figma, reported). Illustrative — re-verify against the Sky1 export before external distribution.

The market pays the meters on disclosure day — not on ship day

Here is the pattern that decides when this thesis pays, and it is not the pattern most investors expect. The market does not pay the meters when they ship an agent. It pays them when the filing names it.

Take the same companies at two different moments. When there was an AI story in the deck and nothing in the filings — Snowflake, JFrog, Cloudflare, quarters earlier — the result was dead money for months, and in one case a 48% drawdown. The proof was public the entire time. The tape did not move.

When agent usage appeared in the filings, backlog and retention confirmed conversion, and guidance was raised on agent products specifically, the same names re-priced roughly +27% on average across seven marked events, overnight or within the print week. Disclosure is the trigger. Everything before disclosure is preparation.

Figure 5. Players · fundamentals · what happened at the print

Meters with receipts

Snowflake · JFrog · Cloudflare

- Agent usage disclosed in filings
- Backlog + retention confirm conversion
- Guidance raised on agent products

~ +27% avg overnight / print-week re-pricing, 7 marked events

Narrative before receipts

Same companies, quarters earlier

- AI story in the deck, nothing in the filings
- Dead money for months — or -48% drawdowns
- Proof public, tape unmoved

~ 0% — SNOW fell 9 weeks after shipping its agent, until the CFO named it

Massive market data (adjusted); company filings. Past re-pricings are not predictive. Cloudflare is an industry exhibit, not held in the aggregate.

Snowflake: sixteen weeks from ship to pay


Snowflake is the cleanest illustration we have, and we would rather show one case in full than seven in summary.

The stock peaked at about \$271 in November 2025 and then fell for twenty-two weeks to a trough near \$135 in early April 2026 — a drawdown of roughly 50%. The company shipped its CoCo agent in February 2026, in the middle of that decline, and the tape did not react at all.

Then on 27 May 2026 the CFO named CoCo as the largest single driver of a raised full-year forecast. The stock gained about 36% overnight and 47% over that week, from \$178 to \$261. By late July it was back to \$268, roughly double the April trough.

Nothing about the business changed on 27 May. The only thing that changed was that the filing said out loud what the product had been doing since February — about sixteen weeks earlier. That gap between truth and price is the window this entire thesis is designed to stand inside.

Figure 6. SNOW month-end close, Jan 2025 – Jul 2026 (\$, split/dividend adjusted)



Month-end adjusted closes per Yahoo Finance, retrieved July 2026; the deck's own exhibit uses Massive weekly adjusted closes, so figures may differ slightly. Peak-to-trough on this series is -50.1% (5 Nov 2025 \$271.26 → 8 Apr 2026 \$135.47 weekly). Week of 27 May 2026: \$178 → \$261, +47%; the overnight move on the print was ~+36%. Past performance is not indicative of future results.

PART FOUR

IV

A Royalty on Machine Labor

the conclusion — the portfolio that owns the meters

“Own the meters; collect the royalty.”

SKYCATCHER

INFINITE SOFTWARE · PART IV

Where we invest: layers 2–5

If infinite software runs on metered rails, the only question left is which rail to own. We divide the stack into six layers and we invest in four of them.

We avoid layer one, agentic applications and workflows, because that is the knife fight — the place where production costs approach zero and entrants arrive without end. We avoid layer six, the cloud and compute substrate, because it is GPU-gated and its economics are set by supply we do not control.

Layers two through five are the investable sweet spot: the tool rails of API, MCP, identity and auth; context, data, analytics and memory; build and verify, meaning repos, CI/CD, evals and tests; and observe, govern and secure — traces, policy, cost. These are the runtime meter, the memory meter, the build meter, and the audit meter. They are the least dependent on GPU availability and the most directly tied to software volume, compliance, and enterprise trust.

The fair question is why not simply own the obvious expressions of this wave — NVIDIA, the hyperscalers, the labs. Our answer is that those trades are either crowded or unavailable. NVIDIA is priced for the buildout by everyone who can read a headline; it is the consensus expression, and consensus is in the price. The hyperscalers own layer six, but at three trillion dollars of market value the agent effect is a rounding error on the multiple. The labs are private, and their layer is the one where competition gives the gains to customers — the price war that drives our entire demand curve is their margin. The meters are where the same wave lands un-modeled, in liquid public companies small enough for it to matter. That is the definition of the gap we hunt.

Figure 1. The portfolio · where we invest

1	Agentic applications & workflows the knife fight	AVOID
2	Tool rails — API · MCP · identity & auth the runtime & checkout meters	OWN
3	Context, data, analytics & memory the context & memory meters	OWN
4	Build & verify — repos, CI/CD, evals, tests the build meter	OWN

5	Observe, govern & secure — traces, policy, cost the audit meter	OWN
6	Cloud & compute substrate GPU-gated	AVOID

What gets a meter into the book

We wrote the admission checklist before we owned a share, which is the only time such a list is worth writing.

A usage-priced core, meaning revenue that meters machine actions rather than human seats. Agent cohorts disclosed in the filings, not in the deck — usage we can audit every quarter. Growing retention and accelerating backlog, so the tape confirms conversion rather than mere adoption. Fortress economics: software gross margins, net cash, Rule of 40. And every link in the chain checkable, from adoption through usage through conversion through auditability.

Where a link cannot be checked, the position takes starter weight and nothing more. Each grade maps to a pre-written weight action we execute inside the quarter, which removes the discretion at exactly the moment discretion becomes expensive.

Figure 2. Admission checklist — written before we owned a share

- **Usage-priced core**
revenue meters machine actions, not human seats
- **Agent cohorts in the filings**
not the deck — disclosed usage, auditable each quarter
- **Growing retention and accelerating backlog**
the tape confirms conversion, not just adoption
- **Fortress economics**
software gross margins, net cash, Rule of 40
- **Every link checkable**
adoption → usage → conversion → auditability, or starter weight only

Skycatcher audit framework. Grades map to pre-written weight actions inside the quarter.

The cheapest this growth has ever been offered

And now the part that makes this a trade rather than an observation. This growth is being offered at the cheapest price it has ever carried.

As of July 2026 the five meters trade at 7.5x EV to revenue — a quarter of their 2021 peak. About 21x forward earnings, which is simply the software tape. Roughly 0.4x growth-adjusted, the cheapest on record for this group. They hold \$14 billion of combined net cash, so none of them depends on financing the buildout. And aggregate gross margin sits at 73%, held intact straight through the agent transition.

Re-accelerating revenue at trough-regime multiples, while backlog and overages compound underneath. We have not been offered this combination before in this group, and we do not expect to be offered it again once the disclosures land.

Figure 3. Aggregate valuation and quality, five public meters, July 2026

7.5x

EV / revenue

a quarter of the 2021 peak

~21x

forward P/E

today's software-tape
multiple

0.4x

growth-adjusted

the cheapest on record for
this group

\$14B

combined net cash

no buildout financing
dependence

73%

aggregate gross margin

held through the agent
transition

Company filings and market data, Skycatcher aggregation, Jul 2026. Multiples approximate.

Multiples compressed while fundamentals compounded

It is worth being precise about how this de-rating happened, because it did not happen for the reason a value investor would normally hope.


EV to next-twelve-months revenue went 30x in FY21, then 12x, 10x, 9x, 8x, and 7.5x by July 2026. Forward earnings multiples compressed from 50x in FY24 to 30x in FY25 to 21x today. FY21 through FY23 carry no meaningful earnings multiple because the group was pre-profitability.

Meanwhile the businesses improved on every axis. Operating margin went from -6% to +18%. Retention returned to 120%. The multiple compressed while the fundamentals compounded, which is a description of sentiment rather than of value.


Figure 4. Aggregate five-meter valuation by fiscal year — EV / NTM revenue and forward P/E (approximate)

APPROXIMATE, BY FISCAL YEAR

EV / NTM REVENUE



FORWARD P/E



Skycatcher estimates from market data and filings; multiples approximate, earnings basis non-GAAP. FY21–23 earnings not meaningful (pre-profitability). Verify against exchange data before external use.

Four honest haircuts still leave \$44B against the street’s \$23B


We built our revenue case by taking four honest haircuts, because a number nobody can attack is a number nobody believes.

- 1. Unit prices fall 30 to 50% a year, and all of it is given back to customers.
- 2. Revenue is booked on discounted multi-year commitments rather than at list price.
- 3. Usage converts to revenue at the measured elasticity of 0.5 to 0.7, not at unity.
- 4. Agent revenue is compute-heavier than seat revenue, and we carry that margin cost.

After all four, aggregate 2029 revenue for the five meters is about \$44 billion against the street’s \$23 billion, with a floor case of \$34 billion. That is roughly a 41% compound growth rate from \$15.9 billion today. The \$21 billion annual wedge between the wave and the fade is the entire trade.

Figure 5. Aggregate 2029 revenue, five public meters (\$B)

2029 REVENUE \$B



Skycatcher model vs. consensus estimates. ~41% CAGR from ~\$15.9B today. Illustrative; see assumptions appendix.

The meters’ margin machine, illustrated

Revenue is the easier half of the argument. The margin machine underneath it is what turns a good revenue outcome into a re-rating.

Indexed to FY21 revenue of 100: revenue is 260 today and about 720 in our 2029 case. Gross margin holds at 72 to 74% throughout. Operating profit goes from –6 to 47 to roughly 280, and operating margin from –6% to 18% to something like 35 to 40%. Profit grows roughly 6x from here.

The important word is “built.” Twenty-four points of operating leverage were assembled before the wave arrived, during the years the multiple was compressing. Incremental drop-through of about 50% is not our assumption — it is what the filings already measure, in a range of 45 to 55%.

Figure 6. Aggregate five-meter P&L, actual and illustrative (\$ indexed, FY21 = 100 revenue)

	FY2021	TODAY	2029 — WAVE CASE
Revenue	100	260	~720
Gross margin	73%	74%	~72%
(–) Opex	(79)	(145)	(~240)
Operating profit	(6)	47	~280
Operating margin	–6%	18%	~35–40%

Aggregate of five public meters per filings (FY21, today); 2029 illustrative at ~50% incremental drop-through (measured 45–55% in filings). Not a forecast of any company’s results.

Today's attach says 0.2–0.3 — that's the entry price, not the ceiling

The obvious objection to all of this is that today's measured attach rate is only 0.2 to 0.3 events per token, well below the 20 to 40 we described. We think that measurement is a lag, not a leak, and the mechanics say why.

Committed contracts burn first, so surge usage draws down prepaid credits and appears in backlog rather than revenue — one to two quarters. Overage only bills above the floors, and then at list — another one to two quarters. Pilots meter lightly while production meters everything, because audit, security, and observability attach only when a workload graduates — two to three quarters. And teams optimise before they let agents run, so efficiency arrives before scale — one to two quarters.

Each of those lags closes on a disclosure date. The attach rate we can measure today is the entry price, not the ceiling. The lag is the opportunity.

Figure 7. Why the measured attach is a lag, not a leak — the surge hits backlog before it hits revenue

1	Committed contracts burn first usage draws down prepaid credits — surge appears in RPO, not revenue	1–2 QTRS
2	Overage bills only above the floors then at list price — the spending nobody budgeted	1–2 QTRS
3	Pilots meter lightly; production meters everything audit, security, observability attach when workloads graduate	2–3 QTRS
4	Efficiency comes before scale teams optimize before they let agents rip	1–2 QTRS

Contract mechanics per usage-platform disclosures. Each lag closes on a disclosure date.

PART FIVE

V

Machine Oracle

the underwriting, the risks, and what to do — our sensing edge and our conclusion

*“Courageous action: to do what no one has done, and prove it in
the real world.”*

LEONARDO DA VINCI

INFINITE SOFTWARE · PART V

What you must believe — and how we check it

Everything in this document reduces to one chain with two bridges in it. If you disagree with us, you disagree at a bridge, and we would rather point at them than let them stay implicit.

The chain runs tokens, to agentic events, to meter revenue: about 27x tokens by 2029 against a floor of about 11x; about 130x events on our case against roughly 43x at the floor; and about 2.8x meter revenue against the street's 1.5x fade.

Bridge one is that tokens proxy agent work — the agentic share rising from about 15% to about 60%, with 20 to 40 billable events per thousand agentic tokens. Bridge two is that events convert to revenue, with elasticity holding at 0.5 to 0.7 as the mix shifts toward premium event tiers. If both hold, 2029 revenue is about \$44 billion and the group re-rates. If elasticity compresses instead, the downside is mediocrity rather than disaster — we drift to consensus and own good businesses at a fair price.

Figure 1. The belief chain — tokens → agentic events → meter revenue, two falsifiable bridges

Tokens • ~27x

by 2029 — our tidal-wave call
(floor ~11x)

Agentic events • ~130x

events, our case — ~43x at the floor

Meter revenue • ~2.8x

our case vs the street's 1.5x fade

Bridge 1

Tokens proxy agent work — agentic token share rises
~15% → ~60%; 20–40 events per 1K agentic tokens

Bridge 2

Events convert to revenue — elasticity holds at 0.5–0.7
as mix shifts to premium event tiers

The street isn't bearish on agents — it has never modeled them


We want to be precise about our disagreement with the street, because it is not the disagreement people assume. The street is not bearish on agents. The street has never modelled them.

Both models start from the same place — about 20% aggregate revenue growth in 2026. Then they separate. The event math produces 25%, 31%, and 36% in the years to 2029, because attach deepens as workloads graduate. Every sell-side model produces 15%, 14%, and 13%, because the default shape of any software model is a fade.

Nobody had to argue against machine labor to arrive at the fade. They only had to leave the template alone.

Figure 2. Aggregate revenue growth, y/y — same starting point, two shapes

AGGREGATE REVENUE GROWTH, Y/Y



Growth paths implied by the aggregate revenue models: consensus fade to ~\$23B vs the event-math path to ~\$44B on our call by 2029. Illustrative.

40x: what history pays for acceleration


Which brings us to the multiple, and to the only piece of this thesis that depends on other people changing their minds.

Today's software tape pays about 21x forward earnings. History pays considerably more for re-accelerated growth: Microsoft's cloud re-acceleration peaked near 35x. NVIDIA in its AI acceleration carried about 45x. ServiceNow at durable 25 to 30% growth has held about 53x. The 2021 SaaS and 1999 regimes paid 65x and above.

We underwrite 40x, which is a PEG of roughly one on re-accelerated growth and the conservative end of every episode above. The 65x cases are the caution rather than the plan; both compressed 60 to 85% afterwards. We would rather be paid for earnings we can count than for a regime we would have to hope for.

Figure 3. Forward P/E paid for re-accelerated growth — precedent episodes vs. today's tape

FORWARD P/E



Multiples approximate: industry & MSFT fwd P/E as of Jul 2026; MSFT 2016–21 re-rating; NVDA FY24–26; NOW 3-yr avg; 2021 SaaS / 1999 per compression studies. Sources: Massive, multiples est., simplified & illustrative.

Price the earnings: every path through 40x beats today

So price the earnings across every combination and see what has to be true to lose money.

Combined equity value in 2029 equals 2029 earnings times a forward multiple plus about \$14 billion of net cash, against roughly \$120 billion today. At today's 21x tape: \$119 billion on the street's fade, \$203 billion on the floor, \$334 billion on the wave. At our 40x band: \$214 billion, \$374 billion, \$624 billion. At the 53x durable-growth precedent: \$279 billion, \$491 billion, \$822 billion.

Eight of the nine cells beat today. The single loser requires the fade to be right and today's trough multiple to persist through 2029 simultaneously. Our case is the 40x-by-wave cell — \$624 billion, +420% — and we carry that figure forward.

Figure 4. Combined equity value, 2029 (\$B) = 2029e earnings × forward P/E + ~\$14B net cash · today: ~\$120B

	STREET FADE — ~\$5.0B	FLOOR — ~\$9.0B	WAVE — ~\$15.3B
21x — today's tape	\$119B (−1%)	\$203B (+69%)	\$334B (+178%)
40x — our band	\$214B (+78%)	\$374B (+212%)	\$624B (+420%)
53x — durable-growth precedent	\$279B (+133%)	\$491B (+309%)	\$822B (+585%)

Skycatcher scenario grid, illustrative. 2029e non-GAAP net margins: fade ~22%, floor ~26%, wave ~35%. P/E bands per previous slide. Approximate.

Our case: ~5x over four years


Held to 2029, our case is roughly 5x over four years, and it is worth separating where that return comes from.

Earnings growth at today's 21x contributes about \$214 billion. The re-rate from 21x to 40x contributes about \$290 billion. Roughly half the return is arithmetic we can already measure in the filings; the other half requires the market to pay a normal multiple for abnormal growth. The consensus drift case returns +17%. The reference floor returns +220%.

There is a hyperboom case — approximately \$1 trillion, about 8.3x from here, at 4,000 quadrillion tokens and 1999-regime multiples. We show it because it is the honest right tail. We do not underwrite it, and nothing in our positioning depends on it.

Figure 5. Combined value of the five meters, 2026–2029e (\$B, Skycatcher scenarios)

COMBINED VALUE OF THE FIVE METERS, \$B



Skycatcher scenarios, illustrative; central case capitalizes the unmodeled wedge at 40x — the conservative end of historical re-acceleration regimes. 1996–99 rails paid 6–13x. Endpoints are from the model; 2027 and 2028 are constant-rate interpolations shown to indicate shape, not a year-by-year forecast.

Why now: five things changed in the last eighteen months

We are asked, reasonably, why now rather than two years ago or two years from now. The answer is that five things became true in the last eighteen months, and every one of them was false before.

Agents now do delegation rather than assistance — 77% of API traffic classifies as automation. The price of agent doing has become trivial, at \$0.40 a million tokens, which makes an always-on agent-day cost a few dollars. Machine work acquired billable units, with agentic SKUs shipping across seven vendors. The wave reached the filings, with retention at 120 and climbing, backlog up 50%, and overage moving from 14 to 24% of revenue. And the multiple did not come with any of it — 21x forward against a 40x precedent, because the street models zero agents.

That last line is the entry condition. Everything else on this list is a reason the business will compound. Only the multiple is a reason to buy it today.

Figure 6. The entry condition — every line below was false eighteen months ago

- | | |
|---|---|
| 1 | Agents do delegation
Today — agents ship; 77% of API traffic is automation |
| 2 | The price of “agent doing” is trivial
Today — \$0.40/M; an always-on agent-day costs a few dollars |
| 3 | Machine work got billable units
Today — billable agent SKUs are shipping: AWUs, Pay Per Use, flex credits |
| 4 | The wave reached the filings
Today — NDR 120 and climbing · RPO +50% · overage 14→24% |
| 5 | The multiple didn’t come with it
Today — 21x fwd P/E vs. 40x precedent; the street models zero agents |

Details in the body of the deck: traffic mix, price, billable units, tape, multiples.

Said simply: the three bets you make

Said as simply as we can say it, there are three bets here and they are sequential rather than independent.

First, the wave builds — tokens about 27x by 2029, where the street's own base case is the 11x floor. Second, the wave shifts from agents asking to agents doing — agentic share from 15% to about 60%, billable events about 130x, because events follow tasks rather than tokens. Third, the wave lands on the meter bill — revenue about 2.8x against the 1.5x fade, roughly \$500 billion of value the street is not modelling, and eight of nine priced paths beating today.

The second is the one to get right. If delegation deepens, the rest takes care of itself.

Figure 7. The whole thesis in three lines — each one graded by Sky1

- | | |
|---|---|
| 1 | The wave builds.
Tokens ~27x by 2029 — our tidal-wave call; the street's own base case is the ~11x floor. |
| 2 | The wave shifts from agents asking to doing.
Agentic share 15% → ~60%; billable events ~130x — events follow tasks, not tokens. |
| 3 | The wave lands on the meter bill.
Revenue ~2.8x vs the 1.5x fade — ~\$500B the street isn't modeling; our case ~5.2x, and 8 of 9 priced paths beat today. |

The strongest case against us

The kill conditions in Part V are the tripwires we can grade quarterly. But the strongest argument against this thesis does not trip any of them, so it deserves its own section.

It is token efficiency. Our chain assumes tokens proxy agent work — bridge one. If models improve enough that the same delegated task takes a tenth of the tokens, work decouples from consumption: agents could take over the economy while token demand, and the events attached to it, grow far more slowly than we forecast. Falling prices we have modelled; falling tokens per task we have not.

Our answer is Jevons again, one level down — every efficiency gain in software history has been spent on doing more rather than consuming less, and the observed intensity band already reflects three years of relentless model improvement. But we hold this answer with less confidence than the rest of the document, and we watch tokens-per-task in our own stack as a leading indicator the filings cannot give us.

Second, self-hosting. The sharpest version comes from engineers: if agents make building software free, they also make operating it cheap — an agent can run your Postgres, rotate your certificates, and triage your pagers, so why rent a meter at all? We concede the premise for some layers, and we deliberately hold the commodity-exposed ones lightly. What cannot be self-hosted away is the part of the bill that is really insurance: identity, audit, and compliance sell a third party to blame, and no enterprise counsel accepts its own agent as the counterparty of record. Self-hosted stacks still run on metered compute and delivery, and every agent an enterprise deploys enlarges the surface that has to be observed and governed by someone whose name is on a contract. The watch item is attach pricing in the observability layer — the first place self-hosting would show up.

Third, the buildout itself. If power and compute supply cannot carry 27x token demand, the wave caps at the floor no matter what elasticity says. Two answers. The floor is priced in our grid — at 21x today's multiple the floor case still returns +69%, and eight of nine paths beat today. And a supply-capped wave concentrates pricing power in whoever meters scarce capacity efficiently, which is not a bad description of the companies we own.

Last, the ones we can only size, not solve. Vertical integration: the hyperscalers could bundle identity, observability, and build into the compute bill, turning the layer we own into a feature — against this we hold the meters with the deepest enterprise moats and watch attach pricing, not just attach volume. And concentration: five names is not a diversified thesis, it is a conviction basket, and we size it as the manifesto sizes everything — deep conviction, pre-written exits, and no pretence that the two are a substitute for each other.

Upcoming inflections

Forecasts that cannot be graded are marketing. Ours are dated, keyed to third-party filings, and scored every quarter.

Between August and November 2026 the prints should confirm the turn: net dollar retention at or above 121%, backlog growth above 45%. In 2027 the seat recession goes visible — a major vendor reporting shrinking seats alongside growing revenue, and the arrival of what we call “agent GAAP.” In 2028 machine-work units go mainstream: billable events past 5 quadrillion a year, overage above 30% of aggregate meter revenue. In 2029 the wave lands in full, with tokens crossing two quintillion a year and the five meters billing about \$44 billion.

These are the dates on which we expect to be told we were right or wrong, by somebody other than ourselves.

Figure 8. The Machine Oracle scoreboard

AUG–NOV 2026 The prints confirm the turn

NDR $\geq 121\%$ · backlog growth $> 45\%$

2027 The seat recession goes visible

shrinking seats, growing revenue at a major vendor · “agent GAAP” arrives

2028 Machine-work units go mainstream

billable events pass 5Q/yr; overage $> 30\%$ of aggregate meter revenue

2029 The wave lands in full

tokens cross 2 quintillion/yr; the five meters bill $\sim \$44\text{B}$

Forecasts are falsifiable by design: dated thresholds keyed to third-party filings, graded quarterly. Forward-looking statements; see disclosures.

How you'll know if we're wrong

And here is how you will know if we are wrong, written in advance and keyed to public filings rather than to our own judgement after the fact.

If token prices refuse to fall for two consecutive quarters, the buildout is breaking and our first assumption fails. If a full 10x price tier down buys only about 10x more demand, elasticity has collapsed to the street's floor and we drift to consensus. If the API-versus-consumer token split stalls for a year, the mix shift from asking to doing has stopped. If overage share reverses, the budget governor is binding for years rather than quarters.

Each trigger maps to a pre-written weight action executed inside the quarter. The discipline runs in the other direction too: growth-adjusted above 1.2x and we sell into the parabola.

Figure 9. Kill conditions — written in advance, keyed to public filings

- **Token prices refuse to fall for 2 consecutive quarters**
the buildout is breaking — assumption 1 fails
- **A full 10x price tier down buys only ~10x demand**
elasticity collapsed to the street's floor — we drift to consensus
- **API-vs-consumer token split stalls for a year**
the mix shift from asking to doing stops
- **Overage share reverses**
the budget governor is binding for years, not quarters

Each trigger maps to a pre-written weight action executed inside the quarter. Multiple discipline: growth-adjusted >1.2x = sell into the parabola.

Waiting for proof misses the disclosure-day gap up

One more observation about timing, because it determines what waiting costs.


Across individual holdings in our aggregate index, disclosure re-pricings have run +36%, +32.7%, +31%, +27%, +23.7%, +21%, and +16% — an average of roughly +27%, delivered overnight or inside the print week.

Treat the average with the respect a seven-event sample deserves — it is a small set, drawn from our own index, and a skeptic should ask how the events were marked before crediting the mean. The direction is the claim, not the decimal: agent disclosures have re-priced these businesses upward, at the print, every time one has landed so far.

There are twelve prints between here and 2029. Capital that is in before the print owns that gap. Capital that waits for proof pays it, because the proof and the re-pricing arrive on the same morning. Past re-pricings are not predictive of future ones, and we are not promising a gap up. We are pointing out that the reward for being early here is not patience — it is a specific, repeatedly observed discount.

Figure 10. Disclosure re-pricings — individual holdings in the aggregate index — overnight / print-week moves

RE-PRICING AT THE PRINT



Overnight close-to-close and print-week moves for individual holdings in the Skycatcher aggregate index, shown anonymized; per Massive market data (adjusted). Average illustrative; past re-pricings are not predictive of future moves.

What to do with this

We have tried to write a document that can be acted on, not admired. If you take the argument, the action follows from the timing math: the re-pricing happens at disclosure, disclosure happens at the prints, and there are twelve prints between here and 2029.

Capital positioned before the August–November 2026 prints owns the first and largest of the gaps. Capital that waits for the filings to confirm will buy the same businesses roughly 27% higher, on the average of the seven re-pricings we have already watched. Waiting for proof is not prudence in this setup; it is paying retail for it.

And if we are wrong, the triggers above say so in public filings, on dates we do not control, with pre-written weight actions attached. That is the deal we are offering: a falsifiable thesis, graded quarterly, at machineoracle.ai.

Five public companies sit at the tollbooths of machine labor. With roughly \$120 billion of combined market value against \$15.9 billion of revenue, they meter the infrastructure every agent must use — every query, transaction, write, test, and log. Thirty-one years ago, an internal memo named the last wave while the street still called it a fad. This document is ours:

**“Machine labor is a tidal wave.
It changes the rules.”**

— SKYCATCHER · JULY 2026



Every scenario in our sensitivity analysis produces upside; the central case points to roughly 5x over four years, and the analogous rails of the internet era returned 5 to 13x from exactly this stage — while the argument was still being had. It is 1996 by the internet clock. The meters are already running. What remains is discipline: the scoreboard is live, the kill conditions are armed, and the next print is weeks away.

Disclosures

What follows is the full disclosure language governing everything above. We would rather you read it than skip it: it names which figures are audited and which are ours, where our estimates carry error bars, which companies we hold and which appear only as exhibits, and where our own interests conflict with a disinterested reading.

Forward-looking statements. This presentation contains forecasts, projections and forward-looking statements, including token-volume scenarios, event-conversion estimates, revenue paths and valuation grids. These are based on assumptions that involve known and unknown risks; actual results may differ materially. Projections are inherently unreliable and should not be relied upon in making an investment decision.

Skycatcher estimates. Aggregate figures for the five-meter basket (retention, RPO, overage share, margins, multiples) are compiled from company filings and restated approximately for basket changes; they have not been audited. Token-volume figures are anchored to platform disclosures and are estimates with error bars. Event-conversion figures derive from Skycatcher's internal model and its own operating stack.

Companies and market data. Company references are informational and are not recommendations. Cloudflare and JFrog exhibits illustrate market behavior; Cloudflare is not held in the aggregate. Market data per Massive (adjusted closes). Historical multiples for 1995–99 rails are approximate from public data. Past performance, including past disclosure re-pricings, is not indicative of future results.

AI risk. AI is an emerging technology subject to a higher level of risk and uncertainty than established sectors, including regulatory, supply-chain, capex-cycle and adoption risks. Chapter-VII-style assumptions and kill triggers are maintained at machineoracle.ai and graded quarterly.

Fund performance. References to prior fund returns (including +160% for the Battle Royale period and +60% for the Game Console App Store period) are overall fund returns for the thesis periods shown, per Skycatcher records, are unaudited, and are presented for illustration of investment process only. Position-level attribution is available in the data room. Individual investor results may differ based on timing and terms. Past performance is not indicative of future results, and there can be no assurance that any fund objective will be achieved.

Basis of presentation. These materials are prepared solely for discussion with a limited number of qualified investors and are confidential. They do not constitute an offer to sell or a solicitation of an offer to buy any security; any such offer will be made only through definitive offering documents, which will contain additional information including risk factors and should be read in their entirety. Recipients should consult their own advisers as to legal, tax, and financial matters. Distribution or reproduction without Skycatcher's consent is prohibited.

Conflicts of interest. Skycatcher, its affiliates and principals hold venture-capital positions in certain private companies referenced in this presentation (including SpacetimeDB, SuperTuned.AI and HeyMido), operate DeadWeb.com, and may benefit from increased awareness of these companies. Venture positions are unrealized, illustrative in valuation, and involve a high risk of loss. Sky1 is an internal research tool; its outputs are not independent research.

Trademarks. All third-party names, logos and trademarks (including Nintendo, Google, Microsoft, Sony, Snowflake, Cloudflare, JFrog, Salesforce, Intercom, Zendesk, Cognition, ServiceNow and others) are the property

of their respective owners. Their use is for identification only and does not imply affiliation, sponsorship or endorsement.